

Convolutional Neural Networks

Yi-Ting Tsai

National Sun Yat-sen University

2018/10/24

Outline

Review-example

Pooling Layer

CNN Architectures

LeNet-5 architecture

AlexNet architectures

GoogLeNet architectures

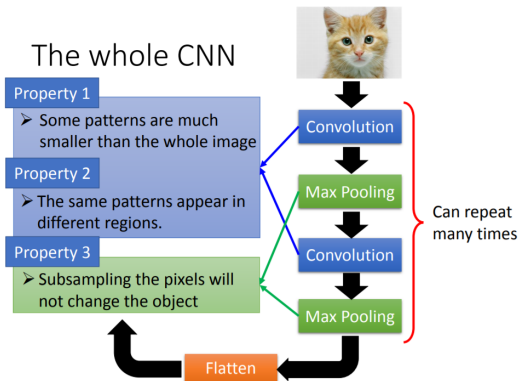
ResNet architectures

SENet architectures

Why CNN for Image:

- Some patterns are much smaller than the whole image
- The same patterns appear in different regions.
- Subsampling the pixels will not change the object

Example:



Example:

CNN – Convolution

Those are the network parameters to be learned.

| | | | | | |
|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 0 | 1 |
| 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 1 | 0 | 0 |
| 1 | 0 | 0 | 0 | 1 | 0 |
| 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 0 | 1 | 0 |

6 x 6 image

| | | |
|----|----|----|
| 1 | -1 | -1 |
| -1 | 1 | -1 |
| -1 | -1 | 1 |

Filter 1

Matrix

| | | |
|----|---|----|
| -1 | 1 | -1 |
| -1 | 1 | -1 |
| -1 | 1 | -1 |

Filter 2

Matrix

⋮

Property 1

Each filter detects a small pattern (3 x 3).

Example:

CNN – Convolution

stride=1

| | | | | | |
|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 0 | 1 |
| 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 1 | 0 | 0 |
| 1 | 0 | 0 | 0 | 1 | 0 |
| 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 0 | 1 | 0 |

6 x 6 image


| | | |
|----|----|----|
| 1 | -1 | -1 |
| -1 | 1 | -1 |
| -1 | -1 | 1 |

Filter 1

3 -1

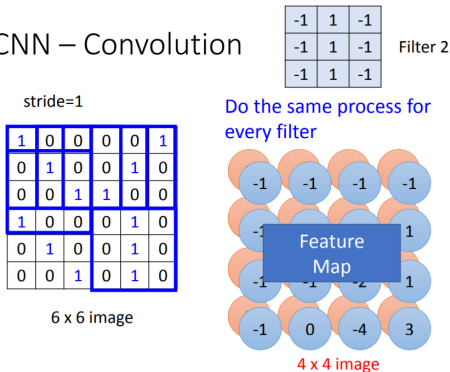
Example:

CNN – Convolution




Example:

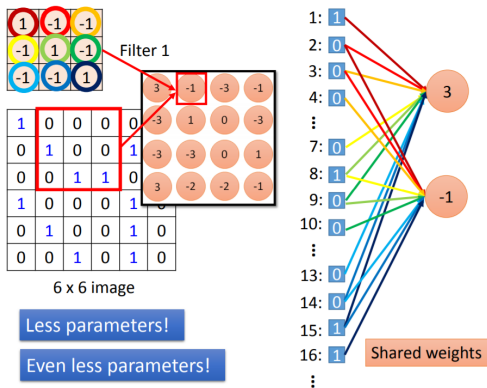
CNN – Convolution



Example:




Example:



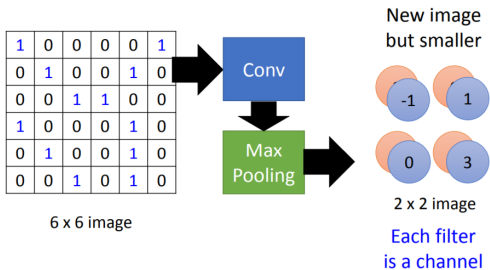
Example:

CNN – Max Pooling



Example:

CNN – Max Pooling



Pooling Layer:

Their goal is to *subsample* (i.e., shrink) the input image in order to reduce the computational load, the memory usage, and the number of parameters




Figure 13-8. Max pooling layer (2×2 pooling kernel, stride 2, no padding)

Figure: Max pooling layer(2×2 pooling kernel, stride 2, no padding)

Typical CNN architectures:




Figure 13-9. Typical CNN architecture

Figure: Typical CNN architecture

CNN Architectures:

A good measure of this progress is the error rate in competitions such as the ILSVRC ImageNet challenge. In this competition the top-5 error rate for image classification fell from over 0.26 to barely over 0.03 in just five years.

We will first look at the classical LeNet-5 architecture (1998), then three of the winners of the ILSVRC challenge: AlexNet (2012), GoogLeNet (2014), ResNet(2015) , and SENet(2017) .

CNN Architectures:




Figure: ILSVRC ImageNet challenge

LeNet-5 architecture:

Table 13-1. LeNet-5 architecture

| Layer | Type | Maps | Size | Kernel size | Stride | Activation |
|-------|-----------------|------|----------------|--------------|--------|------------|
| Out | Fully Connected | — | 10 | — | — | RBF |
| F6 | Fully Connected | — | 84 | — | — | tanh |
| C5 | Convolution | 120 | 1×1 | 5×5 | 1 | tanh |
| S4 | Avg Pooling | 16 | 5×5 | 2×2 | 2 | tanh |
| C3 | Convolution | 16 | 10×10 | 5×5 | 1 | tanh |
| S2 | Avg Pooling | 6 | 14×14 | 2×2 | 2 | tanh |
| C1 | Convolution | 6 | 28×28 | 5×5 | 1 | tanh |
| In | Input | 1 | 32×32 | — | — | — |

Figure: LeNet-5 architecture

AlexNet architectures:

Table 13-2. AlexNet architecture

| Layer | Type | Maps | Size | Kernel size | Stride | Padding | Activation |
|-------|-----------------|---------|------------------|----------------|--------|---------|------------|
| Out | Fully Connected | — | 1,000 | — | — | — | Softmax |
| F9 | Fully Connected | — | 4,096 | — | — | — | ReLU |
| F8 | Fully Connected | — | 4,096 | — | — | — | ReLU |
| C7 | Convolution | 256 | 13×13 | 3×3 | 1 | SAME | ReLU |
| C6 | Convolution | 384 | 13×13 | 3×3 | 1 | SAME | ReLU |
| C5 | Convolution | 384 | 13×13 | 3×3 | 1 | SAME | ReLU |
| S4 | Max Pooling | 256 | 13×13 | 3×3 | 2 | VALID | — |
| C3 | Convolution | 256 | 27×27 | 5×5 | 1 | SAME | ReLU |
| S2 | Max Pooling | 96 | 27×27 | 3×3 | 2 | VALID | — |
| C1 | Convolution | 96 | 55×55 | 11×11 | 4 | SAME | ReLU |
| In | Input | 3 (RGB) | 224×224 | — | — | — | — |

Figure: AlexNet architectures

AlexNet architectures:




Figure: AlexNet architectures

GoogLeNet architectures:

| type | patch size/ stride | output size | depth | #1×1 | #3×3 reduce | #3×3 | #5×5 reduce | #5×5 | pool proj | params | ops |
|----------------|-----------------------|----------------|-------|------|----------------|------|----------------|------|--------------|--------|------|
| convolution | 7×7/2 | 112×112×64 | 1 | | | | | | | 2.7K | 34M |
| max pool | 3×3/2 | 56×56×64 | 0 | | | | | | | | |
| convolution | 3×3/1 | 56×56×192 | 2 | | 64 | 192 | | | | 112K | 360M |
| max pool | 3×3/2 | 28×28×192 | 0 | | | | | | | | |
| inception (3a) | | 28×28×256 | 2 | 64 | 96 | 128 | 16 | 32 | 32 | 159K | 128M |
| inception (3b) | | 28×28×480 | 2 | 128 | 128 | 192 | 32 | 96 | 64 | 380K | 304M |
| max pool | 3×3/2 | 14×14×480 | 0 | | | | | | | | |
| inception (4a) | | 14×14×512 | 2 | 192 | 96 | 208 | 16 | 48 | 64 | 364K | 73M |
| inception (4b) | | 14×14×512 | 2 | 160 | 112 | 224 | 24 | 64 | 64 | 437K | 88M |
| inception (4c) | | 14×14×512 | 2 | 128 | 128 | 256 | 24 | 64 | 64 | 463K | 100M |
| inception (4d) | | 14×14×528 | 2 | 112 | 144 | 288 | 32 | 64 | 64 | 580K | 119M |
| inception (4e) | | 14×14×832 | 2 | 256 | 160 | 320 | 32 | 128 | 128 | 840K | 170M |
| max pool | 3×3/2 | 7×7×832 | 0 | | | | | | | | |
| inception (5a) | | 7×7×832 | 2 | 256 | 160 | 320 | 32 | 128 | 128 | 1072K | 54M |
| inception (5b) | | 7×7×1024 | 2 | 384 | 192 | 384 | 48 | 128 | 128 | 1388K | 71M |
| avg pool | 7×7/1 | 1×1×1024 | 0 | | | | | | | | |
| dropout (40%) | | 1×1×1024 | 0 | | | | | | | | |
| linear | | 1×1×1000 | 1 | | | | | | | 1000K | 1M |
| softmax | | 1×1×1000 | 0 | | | | | | | | |

Table 1: GoogLeNet incarnation of the Inception architecture <https://blog.csdn.net/marsjiao>

Figure: GoogLeNet architectures

GoogLeNet architectures:




Figure: GoogLeNet architectures

ResNet architectures:




Figure 13-12. Residual learning

Figure: ResNet architectures

ResNet architectures:




Figure 13-13. Regular deep neural network (left) and deep residual network (right)

Figure: ResNet architectures

ResNet architectures:

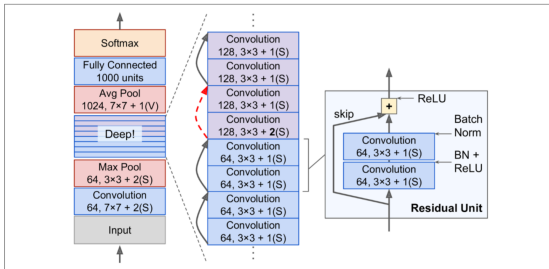


Figure 13-14. ResNet architecture

Figure: ResNet architectures

ResNet architectures:




Figure 13-15. Skip connection when changing feature map size and depth

Figure: ResNet architectures

SENet architectures:

Thanks for listening